

Predicting Customer Spending from Website, Mobile App, and In-Store Activity

June 15, 2020

Annalisa Donat, Data Scientist [Connect with me on LinkedIn](#)

Predicting Customer Spending from Website, Mobile App, and In-Store Activity

Table of Contents

- 1 Background
- 2 Load and Check Data
- 3 Data Transformation
 - 3.1 Email
 - 3.2 Address
 - 3.3 Avatar
- 4 Explore Dataset
 - 4.1 Create Training and Test Sets
 - 4.2 Define Visualization Functions
 - 4.3 Explore Training Data
 - 4.3.1 Categorical Data
 - 4.3.2 Numerical Data
- 5 Predict Yearly Amount Spent
 - 5.1 Define Visualization Functions
 - 5.2 Linear Regression
 - 5.3 RANSAC Linear Regression
 - 5.4 Decision Tree Regressor
 - 5.5 ADABOOST Regression
- 6 Conclusion

1 Background

Client: An e-commerce company based in New York City that sells high-end clothing inspired by traditional textiles of Guam. Customers come in to the store, have (in-store) sessions/meetings with a personal stylist, then they can go home and order either on a mobile app or website for the clothes they want. The company only ships to areas served by the US Postal Service.

Task: Predict the amount a customer will spend each year based on website, mobile app, and in-store activity.

Data: - Avg. Session Length: Average session of in-store style advice sessions - Time on App: Average time spent on App in minutes - Time on Website: Average time spent on Website in minutes - Length of Membership: How many years the customer has been a member - Avatar: Color of the avatar the customer selected on the website - Email: Email address - Address: Shipping address

2 Load and Check Data

Import libraries:

```
[1]: import numpy as np
import pandas as pd
import re
from scipy.stats import f_oneway, pearsonr, spearmanr
from sklearn.ensemble import AdaBoostRegressor
from sklearn.linear_model import LinearRegression, RANSACRegressor
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor

import matplotlib.pyplot as plt
import seaborn as sns
```

Load the data into a pandas DataFrame:

```
[17]: df = pd.read_csv(
    ".../Ecommerce Customers.csv")
```

View the data:

```
[18]: print("Dataset contains %d rows and %d columns" % df.shape)
```

Dataset contains 500 rows and 8 columns

```
[19]: df.head()
```

```
[19]:
```

	Email \		Address	Avatar \
0	mstephenson@fernandez.com		835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet
1	hduke@hotmail.com		4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen
2	pallen@yahoo.com		24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque
3	riverarebecca@gmail.com		1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown
4	mstephens@davidson-herman.com			

4 14023 Rodriguez Passage\nPort Jacobville, PR 3... MediumAquaMarine

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

Check datatype of each column:

```
[20]: df.dtypes
```

```
[20]: Email                object
      Address             object
      Avatar              object
      Avg. Session Length  float64
      Time on App          float64
      Time on Website      float64
      Length of Membership float64
      Yearly Amount Spent  float64
      dtype: object
```

Check for null values:

```
[21]: df.isna().sum()
```

```
[21]: Email                0
      Address             0
      Avatar              0
      Avg. Session Length 0
      Time on App          0
      Time on Website      0
      Length of Membership 0
      Yearly Amount Spent  0
      dtype: int64
```

Number of unique values feature:

```
[22]: df.nunique()
```

```
[22]: Email          500
      Address        500
      Avatar         138
      Avg. Session Length 500
      Time on App     500
      Time on Website  500
      Length of Membership 500
      Yearly Amount Spent 500
      dtype: int64
```

As a sanity check, use the `.describe()` function to make sure the data does not contain any obvious errors (ex. a negative values):

```
[30]: df.describe().round(1)
```

```
[30]:      Avg. Session Length  Time on App  Time on Website  \
count                500.0         500.0         500.0
mean                  33.1           12.1           37.1
std                     1.0            1.0            1.0
min                   29.5            8.5           33.9
25%                   32.3           11.4           36.3
50%                   33.1           12.0           37.1
75%                   33.7           12.8           37.7
max                   36.1           15.1           40.0

      Length of Membership  Yearly Amount Spent
count                500.0         500.0
mean                  3.5         499.3
std                     1.0           79.3
min                    0.3        256.7
25%                    2.9        445.0
50%                    3.5        498.9
75%                    4.1        549.3
max                    6.9        765.5
```

There are no obvious issues.

3 Data Transformation

3.1 Email

Email is a categorical variable with a unique value for each customer. Thus, it is impossible for Email to have any predictive power in its current form. However, natural language processing and text mining techniques might be useful for extracting meaning from the Email feature.

Potentially meaningful information that could be extracted from Email: - email service provider - simplicity of email address (proxy for age): - a combination of the following variables: - popularity of email service provider - when the email service provider was founded - length of email address

(the part before the '@' sign) - whether email address contains any English words and how common those words are. - Very simple email addresses may indicate that the user is over a certain age: ex. there is a user in the dataset with the email address "pallen@yahoo . com". This user was likely old enough to use email when yahoo mail launched in 1997, as allen is a common first and last name and thus this username would be in high demand. - gender - use a library to detect female and male names in the email address. - ex. there is a user in the dataset with the email "riverarebecca@gmail . com". This is likely a woman named Rebecca Rivera. - education and place of employment: - work email addresses or email addresses ending in .edu may contain this information.

For the purpose of simplicity, the only information I will extract from Email will be the email service provider.

```
[23]: df["email_serv"] = df["Email"].apply(lambda x: re.split("@", x)[1])
```

```
[24]: df["email_serv"].nunique()
```

```
[24]: 244
```

```
[25]: pd.DataFrame(df["email_serv"].value_counts()).head()
```

```
[25]:
```

	email_serv
gmail.com	87
hotmail.com	87
yahoo.com	76
jones.com	2
edwards.com	2

```
[26]: df["email_type"] = df["email_serv"].apply(lambda x: "." + re.split("\.", x)[1])
```

```
[27]: df["email_type"].unique()
```

```
[27]: array(['.com', '.biz', '.net', '.info', '.org'], dtype=object)
```

```
[28]: df["email_serv"] = df["email_serv"].apply(lambda x: re.split("\.", x)[0])
df["email_serv"] = df["email_serv"].apply(
    lambda x: x if x in ["gmail", "hotmail", "yahoo"] else "other")
```

```
[29]: df.head()
```

```
[29]:
```

	Email \
0	mstephenson@fernandez.com
1	hduke@hotmail.com
2	pallen@yahoo.com
3	riverarebecca@gmail.com
4	mstephens@davidson-herman.com

Address

Avatar \

0	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet
1	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen
2	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque
3	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown
4	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent	email_serv	email_type
0	587.951054	other	.com
1	392.204933	hotmail	.com
2	487.547505	yahoo	.com
3	581.852344	gmail	.com
4	599.406092	other	.com

3.2 Address

Just like Email, Address is a categorical variable with a unique value for each customer. Thus, it is impossible for Address to have any predictive power in its current form. However, natural language processing and text mining techniques might be useful for extracting meaning from the Address feature.

Extract state/territory/military location from address:

```
[31]: military_addresses = ["%s %s" % (code, area)
                             for code in ["APO", "DPO", "FPO"]
                             for area in ["AA", "AE", "AP"]]

[32]: def get_state(addr):
    states_found = re.findall("[A-Z][A-Z] ", addr)
    if len(states_found) == 1:
        clean_state = states_found[0].replace(" ", "")[:-1]
        return clean_state
    else:
        for military_addr in military_addresses:
            if military_addr in addr:
                return military_addr
        return ""

[33]: df["State"] = df["Address"].apply(lambda x: get_state(x))
```

Note that the following are abbreviations for US territories and Sovereign Nations served by the US Postal Service, not errors:

Country/Territory	Abbreviation
American Samoa	AS
Guam	GU
Northern Mariana	MP
Marshall Islands	MH
Micronesia	FM
Palau	PL
Puerto Rico	PR
Virgin Islands	VI

50 states + 1 District of Columbia + 9 military locations + 8 other countries/territories = 68 “states” in total.

Given that the client sells clothing inspired by the textiles of Guam, people from islands in Oceania and/or living on military bases may represent a disproportionately large number of customers.

```
[34]: df["State"].nunique()
```

```
[34]: 68
```

```
[35]: df["State"].value_counts()
```

```
[35]: DE      13
      MO      13
      SC      13
      OR      12
      VT      12
      ..
      WA       4
      FPO AP   4
      FPO AA   4
      ID       3
      APO AP   2
      Name: State, Length: 68, dtype: int64
```

There are too many distinct values for State for State to be a meaningful feature. However, proximity to New York City may be significant.

```
[38]: f_oneway(df[df["State"]=="NY"]["Yearly Amount Spent"],
              df[df["State"]!="NY"]["Yearly Amount Spent"])
```

```
[38]: F_onewayResult(statistic=0.4962441746981331, pvalue=0.48148435280235435)
```

```
[44]: f_oneway(df[df["State"].isin(["CT", "NJ"])["Yearly Amount Spent"],
              df[~df["State"].isin(["CT", "NJ"])["Yearly Amount Spent"]])
```

```
[44]: F_onewayResult(statistic=0.02254901349693498, pvalue=0.8806965864769287)
```

Proximity to NYC is not significant, but whether a client resides in Guam may be a meaningful feature.

```
[46]: f_oneway(df[df["State"]=="GU"]["Yearly Amount Spent"],
             df[df["State"]!="GU"]["Yearly Amount Spent"])
```

```
[46]: F_onewayResult(statistic=6.994277104849975, pvalue=0.008435256826885338)
```

```
[51]: print(df[df["State"]=="GU"]["Yearly Amount Spent"].mean())
      print(df[df["State"]!="GU"]["Yearly Amount Spent"].mean())
```

```
414.7013601333333
```

```
500.3417226105265
```

```
[47]: df[df["State"]=="GU"].shape
```

```
[47]: (6, 11)
```

Individuals from Guam did spend less. I will reexamine this difference later in the analysis, when I will have a better understanding of potential third variables that could explain this trend.

Group State by region.

```
[74]: region_states = {
      "northeast": ["CT", "ME", "MA", "NH", "NJ", "NY", "PA", "RI", "VT",
                    "DE", "MD", "DC"],
      "south": ["AL", "AR", "FL", "GA", "KY", "LA", "MD", "MS", "NC", "SC", "TN",
                "VA", "WV", "TX", "OK"],
      "west": ["WA", "MT", "OR", "ID", "WY", "CA", "NV", "UT", "CO", "AK",
               "HI", "AZ", "NM"],
      "midwest": ["WI", "IL", "ND", "SD", "NE", "KS", "MI", "IN", "MN",
                  "IA", "MO", "OH"],
      "territories": ["AS", "GU", "MP", "FM", "PR", "VI", "MH", "PW"],
      "military": military_addresses
    }
```

```
[75]: state_region = {state: region for region in
                     region_states.keys() for state in region_states[region]}
```

```
[76]: df["Region"] = df["State"].map(state_region)
```

```
[ ]: df = df.drop(columns=["Address"])
```

3.3 Avatar

```
[63]: df["Avatar"].value_counts()
```



```
[63]: CadetBlue          7
      SlateBlue         7
      GreenYellow       7
      Cyan              7
      Teal              7
      ..
      PaleGreen         1
      PaleTurquoise     1
      Yellow            1
      DeepSkyBlue       1
      LightSlateGray    1
      Name: Avatar, Length: 138, dtype: int64
```

Just like Email and Address, Avatar is a categorical variable with a unique value for each customer. Thus, it is impossible for Avatar to have any predictive power in its current form. However, given that there is no objective way to group these colors into an overarching hierarchy (ex. should pastel blue be grouped with blues or pastels), I am just going to drop the Avatar column from the dataset.

While Avatar is not useful for this analysis, I would recommend that the company investigate whether there is a relationship between Avatar color and the color of future clothing purchases. If so, avatar color could indicate which colors will be in demand next fashion season.

```
[64]: df = df.drop(columns=["Avatar"])
```

```
[65]: df.head()
```

```
[65]:   Avg. Session Length  Time on App  Time on Website  Length of Membership \
0          34.497268      12.655651      39.577668           4.082621
1          31.926272      11.109461      37.268959           2.664034
2          33.000915      11.330278      37.110597           4.104543
3          34.305557      13.717514      36.721283           3.120179
4          33.330673      12.795189      37.536653           4.446308

      Yearly Amount Spent  email_serv  email_type  State      Region
0          587.951054      other      .com      MI      midwest
1          392.204933      hotmail      .com      CA      west
2          487.547505      yahoo      .com      DC      northeast
3          581.852344      gmail      .com      OH      midwest
4          599.406092      other      .com      PR      territories
```

4 Explore Dataset

4.1 Create Training and Test Sets

Creating a test set allows us to evaluate the variance of our model. Variance is the extent to which patterns detected in the training data can be generalized to unseen data (such as the test set).

```
[78]: train_df, test_df = train_test_split(
      df, test_size=0.20, random_state=1)

[79]: continuous_predictors = ["Avg. Session Length", "Time on App",
                              "Time on Website", "Length of Membership"]
      categorical_predictors = ["email_serv", "email_type", "Region"]

[80]: train_X = train_df[continuous_predictors].values
      train_y = train_df["Yearly Amount Spent"].values
      test_X = test_df[continuous_predictors].values
      test_y = test_df["Yearly Amount Spent"].values
```

4.2 Define Visualization Functions

```
[69]: def calculate_pvalues(dfr, method_func=pearsonr):
      """
      dfr : dataframe
      method_func : function (NOT A STRING)
                  - either pearsonr or spearmanr from scipy.stats
      """
      dfcols = pd.DataFrame(columns=dfr.columns)
      pvalues = dfcols.transpose().join(dfcols, how='outer')
      for r in dfr.columns:
          for c in dfr.columns:
              pvalues[r][c] = round(method_func(dfr[r], dfr[c])[1], 2)
      return pvalues

[70]: def get_p_val_mask(dfr, method_func=pearsonr, p_val_cutoff=0.1):
      """
      dfr : dataframe
      method_func : function (NOT A STRING)
                  - either pearsonr or spearmanr from scipy.stats
      p_val_cutoff : float
                  - 0 < p_val_cutoff < 1
      """
      p_vals_mtrx = calculate_pvalues(dfr, method_func=method_func)
      for c in list(p_vals_mtrx):
          p_vals_mtrx[c] = p_vals_mtrx[c].apply(
              lambda x: 0 if x >= 0.1 else 1)
      return p_vals_mtrx
```


4.3 Explore Training Data

4.3.1 Categorical Data

Region

```
[85]: fig, my_ax = plt.subplots(figsize=(10, 4))
sns.swarmplot(x="Region", y="Yearly Amount Spent", data=train_df, ax=my_ax)
my_ax.set_title("Yearly Amount Spent by Region", fontsize=14)
my_ax.set_ylabel("Yearly Amount Spent (USD)")

plt.show()
```



Conduct an ANOVA test to see if there is a significant difference in the Yearly Amount Spent in each Region:

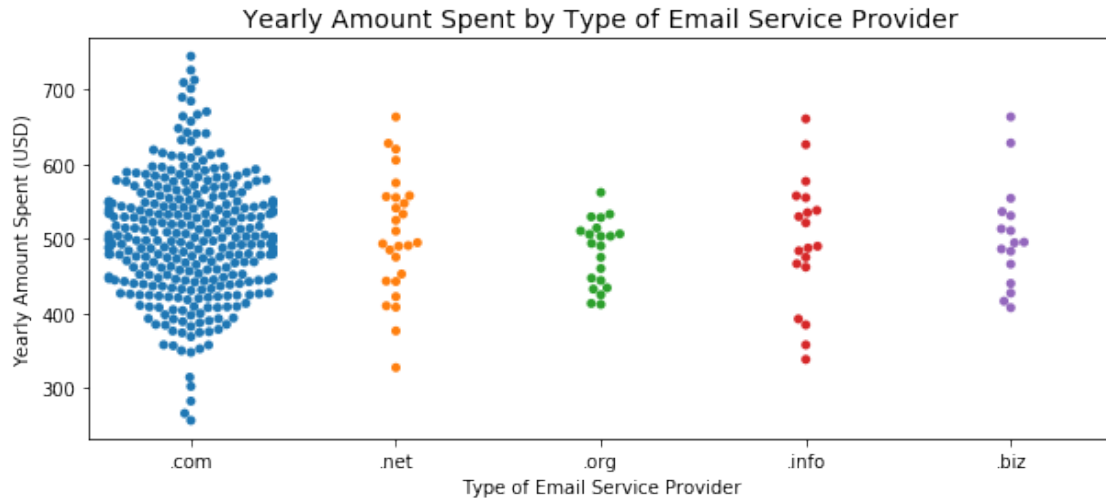
```
[86]: yas_regions = []
for area in train_df["Region"].unique():
    yas_regions.append(train_df[train_df["Region"]==area].copy()["Yearly Amount Spent"])
f_oneway(*yas_regions)
```

```
[86]: F_onewayResult(statistic=0.5507746302861277, pvalue=0.737737378414197)
```

p-value = 0.74 > 0.1, meaning that Region is not a significant predictor of Yearly Amount Spent.

Type of Email Service Provider

```
[88]: fig, my_ax = plt.subplots(figsize=(10, 4))
sns.swarmplot(x="email_type", y="Yearly Amount Spent", data=train_df, ax=my_ax)
my_ax.set_title("Yearly Amount Spent by Type of Email Service Provider",
                fontsize=14)
my_ax.set_ylabel("Yearly Amount Spent (USD)")
my_ax.set_xlabel("Type of Email Service Provider")
plt.show()
```



Conduct an ANOVA test to see if there is a significant difference in the Yearly Amount Spent for each email_type:


```
[89]: yas_email_types = []
      for email_sp_type in train_df["email_type"].unique():
          yas_email_types.append(train_df[train_df["email_type"]==email_sp_type
                                         ].copy()["Yearly Amount Spent"])
      f_oneway(*yas_email_types)
```

```
[89]: F_onewayResult(statistic=0.34281407462547203, pvalue=0.8489876216512878)
```

p-value = 0.74 > 0.1, meaning that email_type is not a significant predictor of Yearly Amount Spent.

Email Service Provider

```
[90]: fig, my_ax = plt.subplots(figsize=(10, 4))
      sns.swarmplot(x="email_serv", y="Yearly Amount Spent", data=train_df, ax=my_ax)
      my_ax.set_title("Yearly Amount Spent by Email Service Provider",
                      fontsize=14)
      my_ax.set_ylabel("Yearly Amount Spent (USD)")
      my_ax.set_xlabel("Email Service Provider")
      plt.show()
```



Conduct an ANOVA test to see if there is a significant difference in the Yearly Amount Spent for each email_serv:

```
[91]: yas_email_serve = []
for email_service_provider in train_df["email_serv"].unique():
    yas_email_serve.append(train_df[train_df["email_serv"]==email_service_provider
                                  ].copy()["Yearly Amount Spent"])
f_oneway(*yas_email_serve)
```

```
[91]: F_onewayResult(statistic=1.0082937956670057, pvalue=0.38898311660136486)
```

p-value = 0.39 > 0.1, meaning that email_serv is not a significant predictor of Yearly Amount Spent.

4.3.2 Numerical Data

```
[92]: continuous_vars = continuous_predictors + ["Yearly Amount Spent"]

fig, axes = plt.subplots(nrows=5, ncols=5, figsize=(12, 12))

# plot data:
for j in range(5):
    for i in range(5):
        if i==j:
            axes[j][i].hist(x=train_df[continuous_vars[i]], bins=20,
                             color="blue", alpha=0.75)
            axes[j][i].set_yticklabels([])
        else:
            axes[j][i].scatter(x=train_df[continuous_vars[i]],
```

```

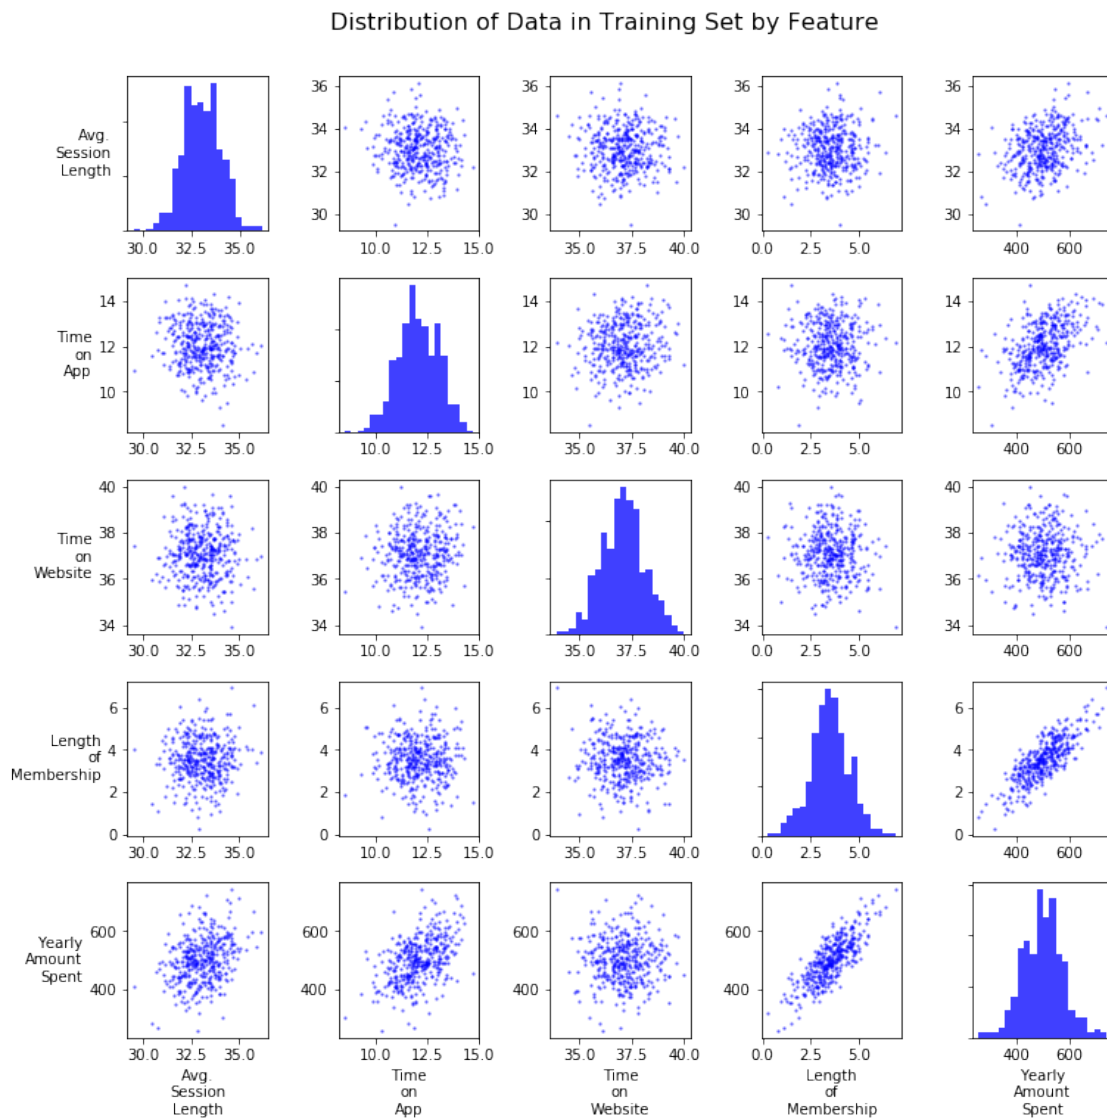
y=train_df[continuous_vars[j]],
color="blue", s=2, alpha=0.4)

# formatting: label x-axis of bottom row and y-axis of left column
for k in range(5):
    label_text = continuous_vars[k].replace(" ", "\n")
    axes[4][k].set_xlabel(label_text)
    axes[k][0].set_ylabel(label_text, rotation=0, ha="right", va="center")

fig.suptitle("Distribution of Data in Training Set by Feature",
            fontsize=16, y=.93)

plt.subplots_adjust(wspace=.5, hspace=.3)
plt.show()

```



All variables have a normal distribution.

```
[93]: fig, axes = plt.subplots(figsize=(9, 3.5), ncols=2)
sns.heatmap(train_df[continuous_vars].corr(
    method="pearson").round(2)*get_p_val_mask(train_df[continuous_vars],
    ↪pearsonr),
    cmap="Blues", ax=axes[0], center=.5, annot=True, cbar=False)
axes[0].set_title("Pearson")

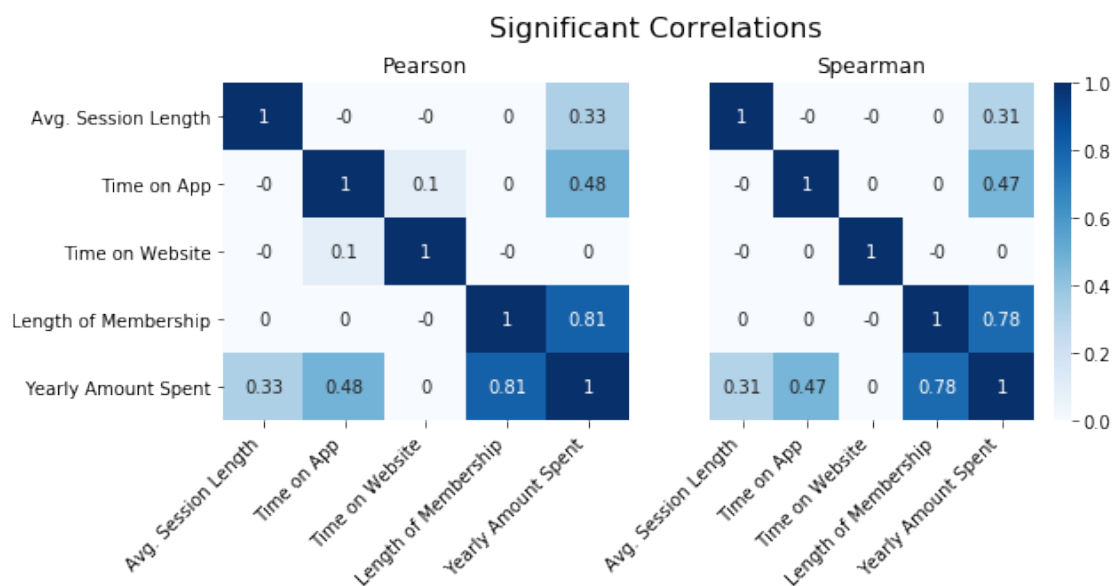
sns.heatmap(train_df[continuous_vars].corr(
    method="spearman").round(2)*get_p_val_mask(train_df[continuous_vars],
    ↪spearmanr),
    cmap="Blues", ax=axes[1], center=.5, annot=True)
axes[1].set_title("Spearman")

axes[1].set_yticklabels([])
axes[1].set_yticks([])

for my_ax in axes:
    bottom, top = my_ax.get_ylim()
    my_ax.set_ylim(bottom + 0.5, top - 0.5)
    my_ax.set_xticklabels(my_ax.get_xticklabels(), rotation=45, ha="right")

fig.suptitle("Significant Correlations", y=1.03, fontsize=16)

plt.show()
```



Observations: Avg. Session Length, Time on App, and Length of Membership all have a positive linear relationship with Yearly Amount Spent and are not correlated with each other. Length of Membership is the greatest predictor of Yearly Amount Spent, followed by Time on App, and Avg. Session Length. There is no relationship between Time on Website and Yearly Amount Spent.

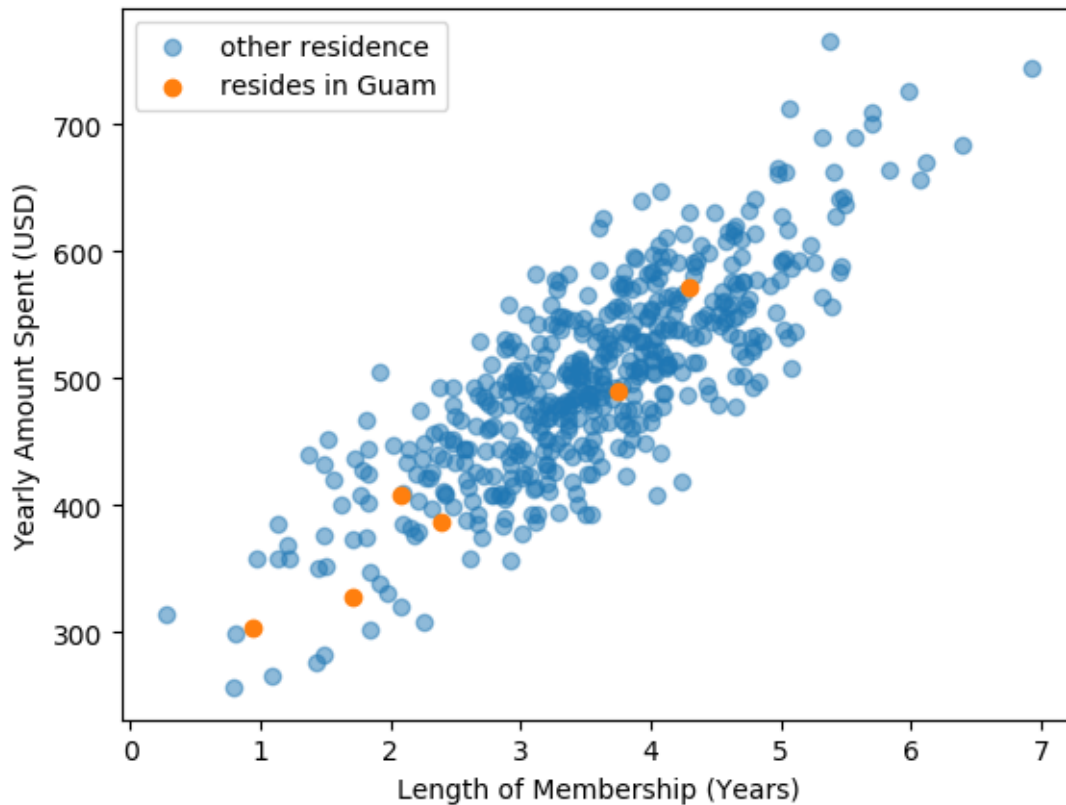
Implications for Model: As the three predictor variables have no relationship with each other, a Principal Components Analysis is not needed. Furthermore, a Multivariate Linear Regression (perhaps with RANSAC) will likely fit the data well. Given the lack of complex relationships in the dataset, more complex methods, such as a Decision Tree Regression or AdaBoost Regression, will not suit the data as well.

Effect of Residence in Guam, Revisited

As previously discussed, customers living in Guam have lower values for Yearly Amount Spent. A quick scatter plot of Length of Membership (the strongest predictor of Yearly Amount Spent) and Yearly Amount Spent will show whether a l

```
[132]: fig, ax = plt.subplots()
ax.scatter("Length of Membership", "Yearly Amount Spent",
           data=df[df["State"]!="GU"], alpha=.5, label="other residence")
ax.scatter("Length of Membership", "Yearly Amount Spent",
           data=df[df["State"]=="GU"], label="resides in Guam")
ax.set_ylabel("Yearly Amount Spent (USD)")
ax.set_xlabel("Length of Membership (Years)")

ax.legend()
plt.show()
```

Given that only 6 customers reside in Guam and that their lower Yearly Amount Spent values are proportional to their lower Length of Membership values, I will not include whether a customer resides in Guam from the model.

5 Predict Yearly Amount Spent

```
[94]: significant_predictors = ["Avg. Session Length",
                              "Time on App",
                              "Length of Membership"]
```

```
[95]: train_X = train_df[significant_predictors].values
      test_X = test_df[significant_predictors].values
```

5.1 Define Visualization Functions

```
[97]: color_explanation = "Darker purple dots correspond to higher values for "\
                          "the sum of the standard deviations of the other two predictor variables."
```

```
[98]: sc = StandardScaler()
```

```
[99]: def plot_model_and_residuals(model, thin_trendline=False, model_name=None):
    plt.rcParams()
    fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(12, 8))

    for i in range(3):
        model = model.fit(train_X[:, i].reshape(-1, 1), train_y)
        x_span = train_X[:, i].max()-train_X[:, i].min()
        x_all = np.linspace(train_X[:, i].min()-int(x_span/6),
                             train_X[:, i].max()+int(x_span/6), 100)
        y_all_pred = model.predict(x_all.reshape(-1, 1))
        other_X_cols = [j for j in range(3) if j!=i]
        ja_std = sc.fit_transform(train_X[:, other_X_cols[0]].reshape(-1, 1))
        jb_std = sc.fit_transform(train_X[:, other_X_cols[1]].reshape(-1, 1))
        hue_col = ja_std + jb_std
        hue_col = hue_col.reshape(train_X[:, i].shape)
        sns.scatterplot(train_X[:, i], train_y,
                        hue=hue_col,
                        alpha=.7, ax=axes[0][i])
        if thin_trendline:
            axes[0][i].plot(x_all, y_all_pred, color="black",
                            linewidth=1, alpha=0.6)
        else:
            axes[0][i].plot(x_all, y_all_pred, color="black")
        axes[0][i].set_title(significant_predictors[i])
        axes[0][i].set_xlabel(significant_predictors[i])
        axes[0][i].set_ylabel("Yearly Amount Spent")
        axes[0][i].legend_.remove()

        pred_y = model.predict(train_X[:, i].reshape(-1, 1))
        #print(train_y.mean())
        #print(pred_y.mean())
        sns.scatterplot(train_y, train_y-pred_y,
                        hue=hue_col,
                        ax=axes[1][i], alpha=0.6)
        axes[1][i].axhline(color="black", linewidth=1)
        axes[1][i].set_xlabel("Yearly Amount Spent")
        axes[1][i].set_ylabel("Residuals")
        axes[1][i].set_ylim(-275, 275)
        axes[1][i].set_title(significant_predictors[i])
        axes[1][i].legend_.remove()

    if model_name:
        fig.suptitle(model_name + ": Univariate Fit and Residuals",
                     fontsize=14)
    else:
        fig.suptitle("Univariate Fit and Residuals", fontsize=14)
```

```
plt.figtext(0.5, 0.01, color_explanation, horizontalalignment='center')
plt.subplots_adjust(hspace=.4, wspace=.4)
plt.show()
```

```
[100]: def evaluate_model(model):
        model = model.fit(train_X, train_y)
        print("Training Set Score: ",
              model.score(train_X, train_y).round(3))
        print("Test Set Score:      ",
              model.score(test_X, test_y).round(3))
```

```
[101]: def summarize_gs(gs):
        """Prints best test (validation) score of a fitted GridSearchCV object
        Parameters
        -----
        gs : fitted GridSearchCV object

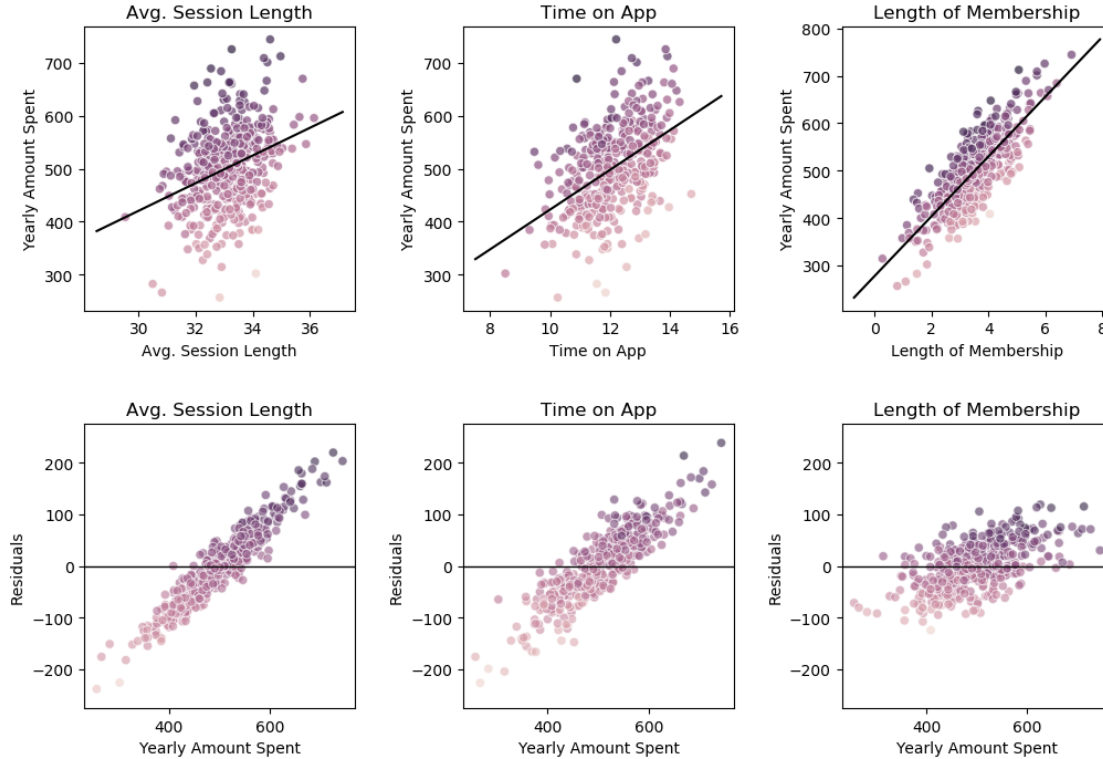
        Returns
        -----
        None
        """
        if len(list(gs.best_params_.keys())) > 0:
            print("Best Parameters:", gs.best_params_)
        gs_std_test_score = gs.cv_results_["std_test_score"][gs.best_index_]
        print("Best Score: %.3f +/- %.3f" % (gs.best_score_, gs_std_test_score))
```

5.2 Linear Regression

```
[102]: lnr = LinearRegression(normalize=True)
```

```
[103]: pipe_lnr = Pipeline([("sc", StandardScaler()), ("lnr", lnr)])
        plot_model_and_residuals(pipe_lnr)
```

Univariate Fit and Residuals



Darker purple dots correspond to higher values for the sum of the standard deviations of the other two predictor variables.

```
[104]: gs_lnr = GridSearchCV(
    estimator=pipe_lnr, param_grid={},
    scoring="r2", return_train_score=True, cv=5, n_jobs=-1, iid=True)
gs_lnr = gs_lnr.fit(train_X, train_y)
summarize_gs(gs_lnr)
```

Best Score: 0.982 +/- 0.005

```
[105]: evaluate_model(pipe_lnr)
```

Training Set Score: 0.983

Test Set Score: 0.989

Wow! Just as I predicted, the multivariate linear regression model fits the data very well.

5.3 RANSAC Linear Regression

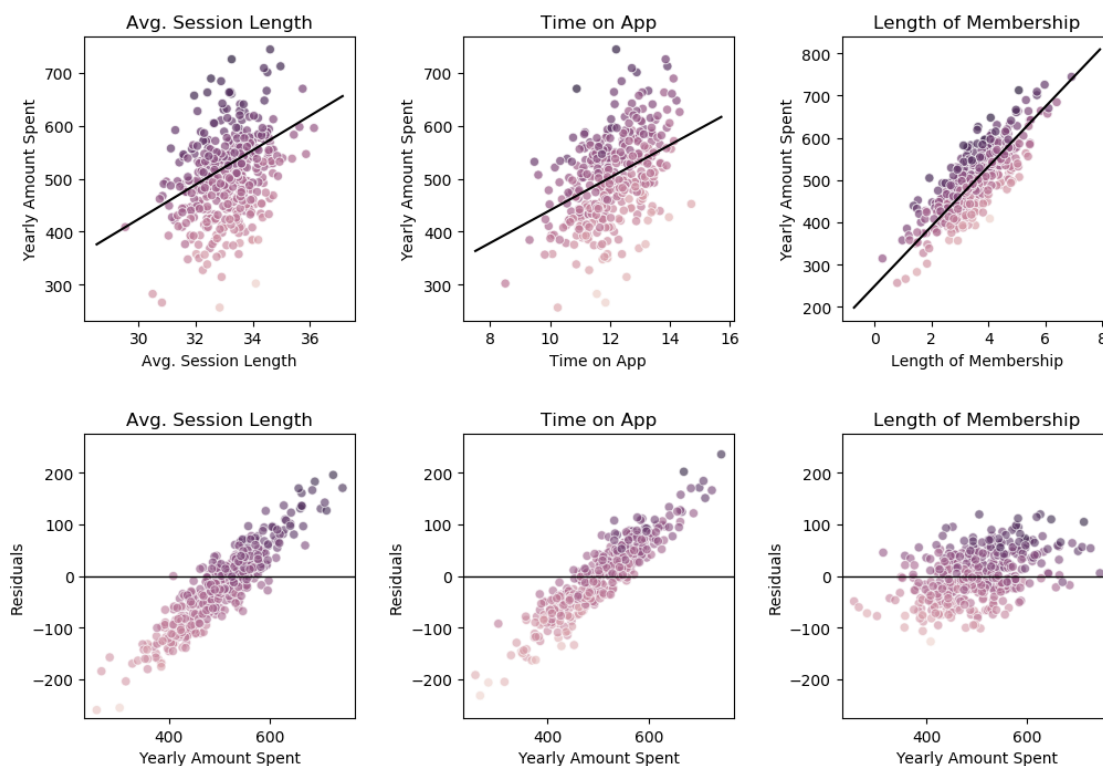
RANSAC (Random Sample Consensus) reduces the influence of outliers in a linear regression. While there is little room for improvement, I will try applying RANSAC to the multivariate linear regression model.

```
[107]: ransac = RANSACRegressor(
        LinearRegression(), max_trials=100, min_samples=50,
        loss='absolute_loss', residual_threshold=3.0, random_state=1)
```

```
[111]: ransac = RANSACRegressor(
        LinearRegression(), max_trials=100,
        loss='absolute_loss', random_state=1)
```

```
[112]: pipe_lnr_ransac = Pipeline([("sc", StandardScaler()), ("ransac", ransac)])
        plot_model_and_residuals(pipe_lnr_ransac,
                                model_name="RANSAC Linear Regression")
```

RANSAC Linear Regression: Univariate Fit and Residuals



Darker purple dots correspond to higher values for the sum of the standard deviations of the other two predictor variables.

Fit multivariate RANSAC linear regression:

```
[113]: gs_lnr_ransac = GridSearchCV(
        estimator=pipe_lnr_ransac,
        param_grid={"ransac__min_samples": [50, 200, 300, 310],
                    "ransac__residual_threshold": [1, 2, 3, 4, 5, 6]
                  },
        scoring="r2", return_train_score=True, cv=5,
```

```
n_jobs=-1, iid=True)

gs_lnr_ransac = gs_lnr_ransac.fit(train_X, train_y)
summarize_gs(gs_lnr_ransac)
```

Best Parameters: {'ransac__min_samples': 300, 'ransac__residual_threshold': 3}
Best Score: 0.982 +/- 0.005

Further fine tune the hyperparameters:

```
[116]: gs_lnr_ransac = GridSearchCV(
        estimator=pipe_lnr_ransac,
        param_grid={"ransac__min_samples": [250, 300, 310],
                    "ransac__residual_threshold": [2.5, 3, 3.5]},
        scoring="r2", return_train_score=True, cv=5,
        n_jobs=-1, iid=True)

gs_lnr_ransac = gs_lnr_ransac.fit(train_X, train_y)
summarize_gs(gs_lnr_ransac)
```

Best Parameters: {'ransac__min_samples': 300, 'ransac__residual_threshold': 2.5}
Best Score: 0.982 +/- 0.005

```
[117]: pipe_lnr_ransac = pipe_lnr_ransac.set_params(**gs_lnr_ransac.best_params_)
```

```
[118]: evaluate_model(pipe_lnr_ransac)
```

Training Set Score: 0.983
Test Set Score: 0.99

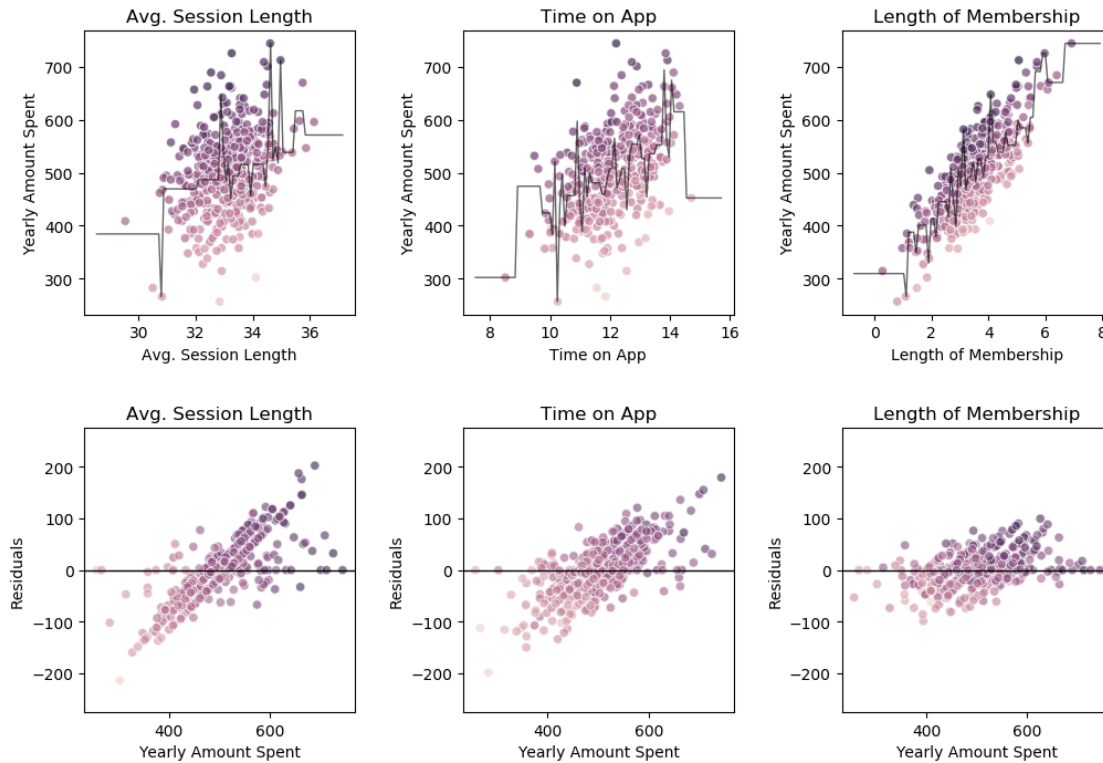
The r2 score for the test set improved by 0.001, but there was very little room for improvement.

5.4 Decision Tree Regressor

As discussed earlier, I do not believe that a Decision Tree Regression will fit the dataset well.

```
[119]: dt_reg = DecisionTreeRegressor(max_depth=8, min_samples_split=4)
        plot_model_and_residuals(dt_reg, thin_trendline=True,
                                model_name="DecisionTreeRegressor")
```

DecisionTreeRegressor: Univariate Fit and Residuals



Darker purple dots correspond to higher values for the sum of the standard deviations of the other two predictor variables.

```
[120]: dtr_params = {"max_depth": [3, 6, 8, 10, 15, 50],
                    "min_samples_split": [2, 5, 20, 30, 50]}
```

```
[121]: gs_dtr = GridSearchCV(
    estimator=DecisionTreeRegressor(random_state=1),
    param_grid=dtr_params,
    scoring="r2", return_train_score=True, cv=5,
    n_jobs=-1, iid=True)
gs_dtr = gs_dtr.fit(train_X, train_y)
summarize_gs(gs_dtr)
```

```
Best Parameters: {'max_depth': 10, 'min_samples_split': 2}
Best Score: 0.833 +/- 0.016
```

```
[122]: dtr = DecisionTreeRegressor(
    random_state=1).set_params(**gs_dtr.best_params_)
```

```
[123]: evaluate_model(dtr)
```

```
Training Set Score: 0.998
```

Test Set Score: 0.914

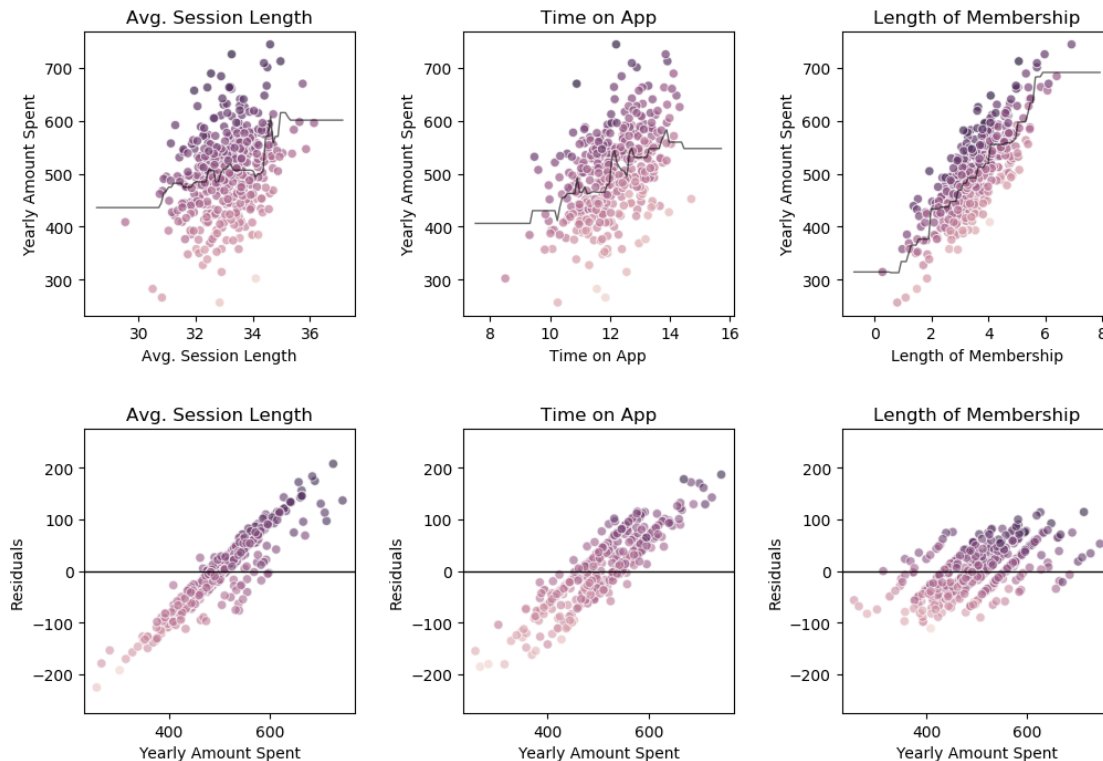
Indeed, the Decision Tree Regression is not as good a model as the Multivariate Linear Regression Model.

5.5 ADABOOST Regression

As discussed earlier, I do not believe that an AdaBoost Regression will fit the dataset well. As the Decision Tree Regression didn't fit the dataset nearly as well as did the Multivariate Linear Regression, I am quite confident that the AdaBoost Regression will not fit the data as well as the Multivariate Linear Regression.

```
[124]: adabr = AdaBoostRegressor(random_state=1)
plot_model_and_residuals(adabr, thin_trendline=True,
                        model_name="AdaBoost Regressor")
```

AdaBoost Regressor: Univariate Fit and Residuals



Darker purple dots correspond to higher values for the sum of the standard deviations of the other two predictor variables.

```
[125]: evaluate_model(adabr)
```

Training Set Score: 0.918

Test Set Score: 0.861

Indeed, the Decision Tree Regression is not as good a model as the Multivariate Linear Regression Model.

6 Conclusion

Avg. Session Length, Time on App, and Length of Membership all have a positive linear relationship with Yearly Amount Spent and are not correlated with each other. Length of Membership is the greatest predictor of Yearly Amount Spent, followed by Time on App, and Avg. Session Length. As the three predictor variables have no relationship with each other, the Multivariate Linear Regression with RANSAC fit the data nearly perfectly. The r^2 score for the test data was 0.990.